

SE183091_NguyenThanhHoa-labassignment

Source Visual Outline

```
1 ---
2 title: "SE183091_NguyenThanhHoa-labassignment"
3 output: html_document
4 date: "2024-10-16"
5 ---
6 ```{r}
7 # Load tidyverse
8 library(tidyverse)
9 # Install required packages
10 install.packages("readr")
11
12 # Load the necessary libraries
13 library(readr)
14 library(httr)
15 ```
```

tidyverse 2.0.0

— Attaching core tidyverse packages —

✓ dplyr	1.1.4	✓ readr	2.1.5
✓ forcats	1.0.0	✓ stringr	1.5.1
✓ ggplot2	3.5.1	✓ tibble	3.2.1
✓ lubridate	1.9.3	✓ tidyr	1.3.1
✓ purrr	1.0.2	— Conflicts —	

tidyverse_conflicts() —

✗ dplyr::filter() masks stats::filter()

✗ dplyr::lag() masks stats::lag()

i Use the conflicted package to force all conflicts to become errors

Error in install.packages : Updating loaded packages

16

```
Source Visual Outline
17 {r}
18
19 # URL to download the file
20 url <- "https://dax-cdn.cdn.appdomain.cloud/dax-airline/1.0.1/lax_to_jfk.tar.gz"
21
22 # Download the file using http
23 GET(url, write_disk("lax_to_jfk.tar.gz", overwrite = TRUE))
24
25
26 {r}
27 # Untar the file in Kaggle
28 untar("lax_to_jfk.tar.gz")
29
30 # Read the CSV file using readr
31 sub_airline <- read_csv("D:/FPT/Kihoc/fall2024/DRS301m/lab/lab3/lax_to_jfk/lax_to_
jfk.csv", col_types = cols(DivDistance = col_number(), DivArrDelay =
col_number()))
32
33 # Check the first few rows
34 head(sub_airline)
35
```

Response [https://dax-cdn.cdn.appdomain.cloud/dax-airline/1.0.1/lax_to_jfk.tar.gz]
Date: 2024-10-16 13:06
Status: 200
Content-Type: application/x-gzip
Size: 58.4 kB
<ON DISK> D:\FPT\Kihoc\fall2024\DRS301m\lab\lab3\lax_to_jfk.tar.gz

A tibble: 6 × 21

Month <dbl>	DayOfWeek <dbl>	FlightDate <date>	Reporting_Airline <chr>	Origin <chr>	Dest <chr>	CRSDepTime <chr>
3	5	2003-03-28	UA	LAX	JFK	2210
11	4	2018-11-29	AS	LAX	JFK	1045
8	5	2015-08-28	UA	LAX	JFK	0805
4	7	2003-04-20	DL	LAX	JFK	2205
11	3	2005-11-30	UA	LAX	JFK	0840
4	1	1992-04-06	UA	LAX	JFK	1450

6 rows | 1-7 of 21 columns

```
36 # Task 1: Data Inspection and Handling Missing Values
37 {r}
38 # 1. Inspect for missing values
39 print("Missing values summary:")
40 print(summary(is.na(sub_airline)))
41 print(sapply(sub_airline, function(x) sum(is.na(x)))) # Missing values per column
42
43
```

[1] "Missing values summary:"

Month	DayOfWeek	FlightDate	Reporting_Airline
Mode :logical	Mode :logical	Mode :logical	Mode :logical
FALSE:2855	FALSE:2855	FALSE:2855	FALSE:2855

Origin	Dest	CRSDepTime	CRSArrTime
Mode :logical	Mode :logical	Mode :logical	Mode :logical
FALSE:2855	FALSE:2855	FALSE:2855	FALSE:2855

DepTime	ArrTime	ArrDelay	ArrDelayMinutes
Mode :logical	Mode :logical	Mode :logical	Mode :logical
FALSE:2855	FALSE:2855	FALSE:2855	FALSE:2855

CarrierDelay	WeatherDelay	NASDelay	SecurityDelay
Mode :logical	Mode :logical	Mode :logical	Mode :logical
FALSE:369	FALSE:369	FALSE:369	FALSE:369
TRUE :2486	TRUE :2486	TRUE :2486	TRUE :2486

LateAircraftDelay	DepDelay	DepDelayMinutes	DivDistance
Mode :logical	Mode :logical	Mode :logical	Mode:logical
FALSE:369	FALSE:2855	FALSE:2855	TRUE:2855

TRUE :2486
DivArrDelay
Mode:logical
TRUE:2855

Month	DayOfWeek	FlightDate	Reporting_Airline
0	0	0	0
Origin	Dest	CRSDepTime	CRSArrTime
0	0	0	0
DepTime	ArrTime	ArrDelay	ArrDelayMinutes
0	0	0	0
CarrierDelay	WeatherDelay	NASDelay	SecurityDelay
2486	2486	2486	2486
LateAircraftDelay	DepDelay	DepDelayMinutes	DivDistance
2486	0	0	2855
DivArrDelay			
2855			

SourceVisual

Outline

```
44
45 ```{r}
46 library(dplyr)
47 library(tidyr)
48 # 2. Handling missing values in specific columns
49 # Replace and remove
50 sub_airline_cleaned <- sub_airline %>%
51   mutate(across(c(CarrierDelay, WeatherDelay, NASDelay, SecurityDelay,
52     LateAircraftDelay), ~replace_na(., 0))) %>%
53   select(-DivDistance, -DivArrDelay)
54 # Check the first few rows of the cleaned data
55 head(sub_airline_cleaned)
56 ```
```

A tibble: 6 × 19

Month <dbl>	DayOfWeek <dbl>	FlightDate <date>	Reporting_Airline <chr>	Origin <chr>	Dest <chr>	CRSDepTime <chr>
3	5	2003-03-28	UA	LAX	JFK	2210
11	4	2018-11-29	AS	LAX	JFK	1045
8	5	2015-08-28	UA	LAX	JFK	0805
4	7	2003-04-20	DL	LAX	JFK	2205
11	3	2005-11-30	UA	LAX	JFK	0840
4	1	1992-04-06	UA	LAX	JFK	1450

6 rows | 1-7 of 19 columns

Source Visual Outline

57
58 {r}
59 print("Missing values summary after handling structural NAs:")
60 print(summary(is.na(sub_airline_cleaned)))
61 print(sapply(sub_airline_cleaned, function(x) sum(is.na(x))))
62

[1] "Missing values summary after handling structural NAs:"

Month	DayOfWeek	FlightDate	Reporting_Airline
Mode :logical	Mode :logical	Mode :logical	Mode :logical
FALSE:2855	FALSE:2855	FALSE:2855	FALSE:2855
Origin	Dest	CRSDepTime	CRSArrTime
Mode :logical	Mode :logical	Mode :logical	Mode :logical
FALSE:2855	FALSE:2855	FALSE:2855	FALSE:2855
DepTime	ArrTime	ArrDelay	ArrDelayMinutes
Mode :logical	Mode :logical	Mode :logical	Mode :logical
FALSE:2855	FALSE:2855	FALSE:2855	FALSE:2855
CarrierDelay	WeatherDelay	NASDelay	SecurityDelay
Mode :logical	Mode :logical	Mode :logical	Mode :logical
FALSE:2855	FALSE:2855	FALSE:2855	FALSE:2855
LateAircraftDelay	DepDelay	DepDelayMinutes	
Mode :logical	Mode :logical	Mode :logical	
FALSE:2855	FALSE:2855	FALSE:2855	
Month	DayOfWeek	FlightDate	Reporting_Airline
0	0	0	0
Origin	Dest	CRSDepTime	CRSArrTime
0	0	0	0
DepTime	ArrTime	ArrDelay	ArrDelayMinutes
0	0	0	0
CarrierDelay	WeatherDelay	NASDelay	SecurityDelay
0	0	0	0
LateAircraftDelay	DepDelay	DepDelayMinutes	
0	0	0	

```
64 1. Columns with the Most Missing Values (Initially):
65
66 The initial inspection revealed that DivDistance, DivArrDelay, CarrierDelay,
67 WeatherDelay, NASDelay, SecurityDelay, and LateAircraftDelay had the most missing
   values.
68 DivDistance and DivArrDelay were missing in all rows (2855 missing values each).
69 The delay-related columns (CarrierDelay, etc.) had 2486 missing values each.
70
71 2. Pros and Cons of Missing Value Strategies:
72
73 Replacing NAs with 0 (for CarrierDelay, WeatherDelay, etc.):
74
75 Pros: This is the most appropriate strategy for these columns. The missingness is
   structural - a missing value indicates a delay of 0, not truly missing data. This
   preserves the information that there was no delay of that particular type.
76
77 Cons: If there were truly missing values mixed in (e.g., data entry errors), this
   method would mask them. However, given the large number of missing values and the
   nature of these variables, this is less likely.
78
79 Removing columns DivDistance and DivArrDelay:
80
81 Pros: Necessary because these columns provided no information (100% missing
   values). Removing them simplifies the dataset and prevents errors in subsequent
   analysis.
82
83 Cons: Loss of potential information if these columns had been populated. However,
   in this case, the loss is unavoidable.
84
85 3. Changes in Dataset Dimensions:
86
87 Original dataset: 2855 rows x 21 columns
88
89 After handling: 2855 rows x 19 columns
90
```

```

91 ▾ ## TASK 2
92 ▾ ```{r}
93 # Simple scaling (dividing by max)
94 sub_airline_cleaned <- sub_airline_cleaned %>%
95   mutate(ArrDelay_simple_scaled = ArrDelay / max(ArrDelay, na.rm = TRUE),
96          DepDelay_simple_scaled = DepDelay / max(DepDelay, na.rm = TRUE))
97 ▴ ```
98
99 ▾ ```{r}
100 # Min-Max Scaling
101 sub_airline_cleaned <- sub_airline_cleaned %>%
102   mutate(ArrDelay_minmax_scaled = (ArrDelay - min(ArrDelay, na.rm = TRUE)) /
103          (max(ArrDelay, na.rm = TRUE) - min(ArrDelay,
104          na.rm = TRUE)),
105          DepDelay_minmax_scaled = (DepDelay - min(DepDelay, na.rm = TRUE)) /
106          (max(DepDelay, na.rm = TRUE) - min(DepDelay,
107          na.rm = TRUE)))
108
109 ▾ ```{r}
110 # Z-score Standardization
111 sub_airline_cleaned <- sub_airline_cleaned %>%
112   mutate(ArrDelay_zscore = (ArrDelay - mean(ArrDelay, na.rm = TRUE)) /
113          sd(ArrDelay, na.rm = TRUE),
114          DepDelay_zscore = (DepDelay - mean(DepDelay, na.rm = TRUE)) /
115          sd(DepDelay, na.rm = TRUE))
116
117 ▾ ```{r}
118 # 3. Compare the results using histograms
119 install.packages("ggplot2")
120 library(ggplot2)
121

```


```

122 {r}
123 # ArrDelay comparison
124 p1 <- ggplot(sub_airline_cleaned, aes(x = ArrDelay_simple_scaled)) +
125   geom_histogram(bins = 30, fill = "blue", alpha = 0.6) +
126   ggtitle("ArrDelay: Simple Scaling") +
127   theme_minimal()
128
129 p2 <- ggplot(sub_airline_cleaned, aes(x = ArrDelay_minmax_scaled)) +
130   geom_histogram(bins = 30, fill = "green", alpha = 0.6) +
131   ggtitle("ArrDelay: Min-Max Scaling") +
132   theme_minimal()
133
134 p3 <- ggplot(sub_airline_cleaned, aes(x = ArrDelay_zscore)) +
135   geom_histogram(bins = 30, fill = "red", alpha = 0.6) +
136   ggtitle("ArrDelay: Z-Score Standardization") +
137   theme_minimal()
138 {r}
139
140 {r}
141 # DepDelay comparison
142 p4 <- ggplot(sub_airline_cleaned, aes(x = DepDelay_simple_scaled)) +
143   geom_histogram(bins = 30, fill = "blue", alpha = 0.6) +
144   ggtitle("DepDelay: Simple Scaling") +
145   theme_minimal()
146
147 p5 <- ggplot(sub_airline_cleaned, aes(x = DepDelay_minmax_scaled)) +
148   geom_histogram(bins = 30, fill = "green", alpha = 0.6) +
149   ggtitle("DepDelay: Min-Max Scaling") +
150   theme_minimal()
151
152 p6 <- ggplot(sub_airline_cleaned, aes(x = DepDelay_zscore)) +
153   geom_histogram(bins = 30, fill = "red", alpha = 0.6) +
154   ggtitle("DepDelay: Z-Score Standardization") +
155   theme_minimal()
156 {r}


```




```
158 {r}  
159 print(p1)  
160  
161
```



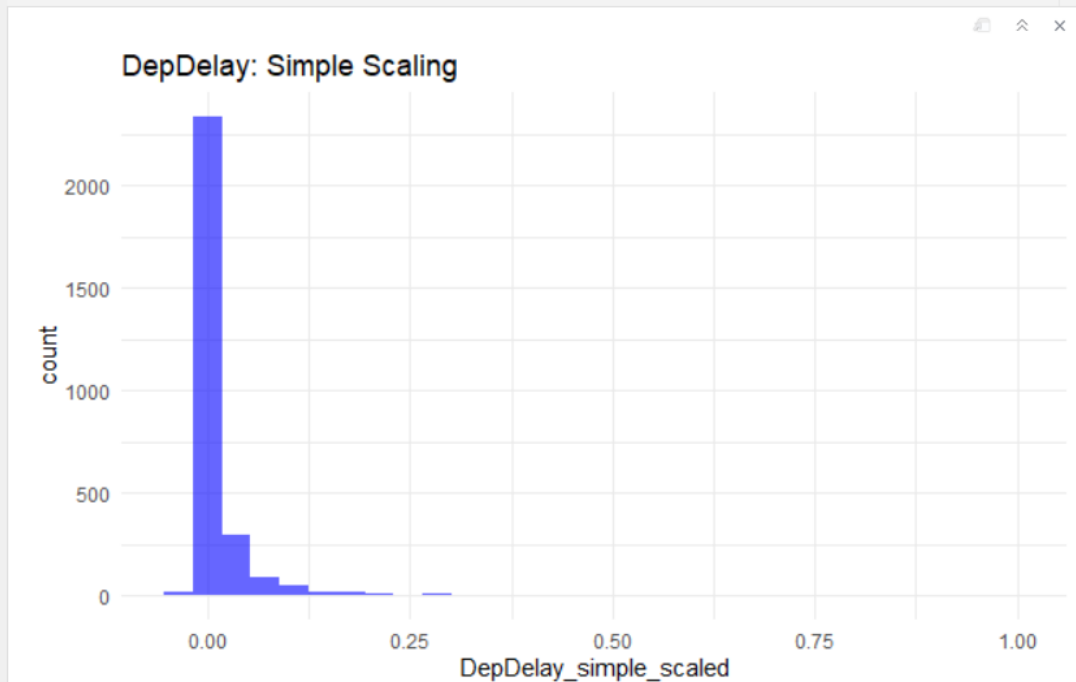
```
163 {r}  
164 print(p2)  
165  
166
```




```
168 {r}  
169 print(p3)  
170
```




```
171 {r}  
172 print(p4)  
173
```



```
174 {r}  
175 print(p5)  
176
```



```
177 {r}  
178 print(p6)  
179
```



180 + Distribution Changes:

181
182 Simple Scaling: The data is rescaled between 0 and 1 based on the maximum value. The shape of the distribution remains the same, but it is compressed into a smaller range.

183
184 Min-Max Scaling: Similar to simple scaling, but ensures the minimum value is 0 and the maximum is 1. The distribution is stretched or squeezed into the 0-1 range.

185
186 Z-score Standardization: Centers the data around 0 with a standard deviation of 1. It transforms the distribution to show how far each value is from the mean, but the shape of the distribution is preserved.

187
188 + Most Appropriate Normalization:

189
190 Z-score Standardization is often the most appropriate when data has outliers, as it scales based on the mean and standard deviation, making it less sensitive to extreme values.

191
192 Min-Max Scaling can be useful if you need all values to be between 0 and 1, but it can be skewed by outliers.

193
194 Simple Scaling may be less suitable because it's entirely dependent on the maximum value, which can distort the scaling if outliers exist.

195
196 + Impact on Further Analysis:

197
198 Z-score Standardization is generally better for statistical analyses like regression or machine learning, where the data needs to be centered and comparable.


199
200 Min-Max Scaling is more useful for algorithms that rely on specific ranges (e.g., neural networks) but may distort distances in the presence of outliers.

201
202 Simple Scaling might not work well for more complex analysis since it doesn't take data distribution into account.

```

203 #TASK 3
204 ```{r}
205 # 1. Create a histogram of the ArrDelay column
206 p_hist <- ggplot(sub_airline_cleaned, aes(x = ArrDelay)) +
207   geom_histogram(bins = 30, fill = "blue", alpha = 0.6) +
208   ggtitle("Histogram of ArrDelay") +
209   theme_minimal()
210
211 # Print the histogram
212 print(p_hist)
213 ```

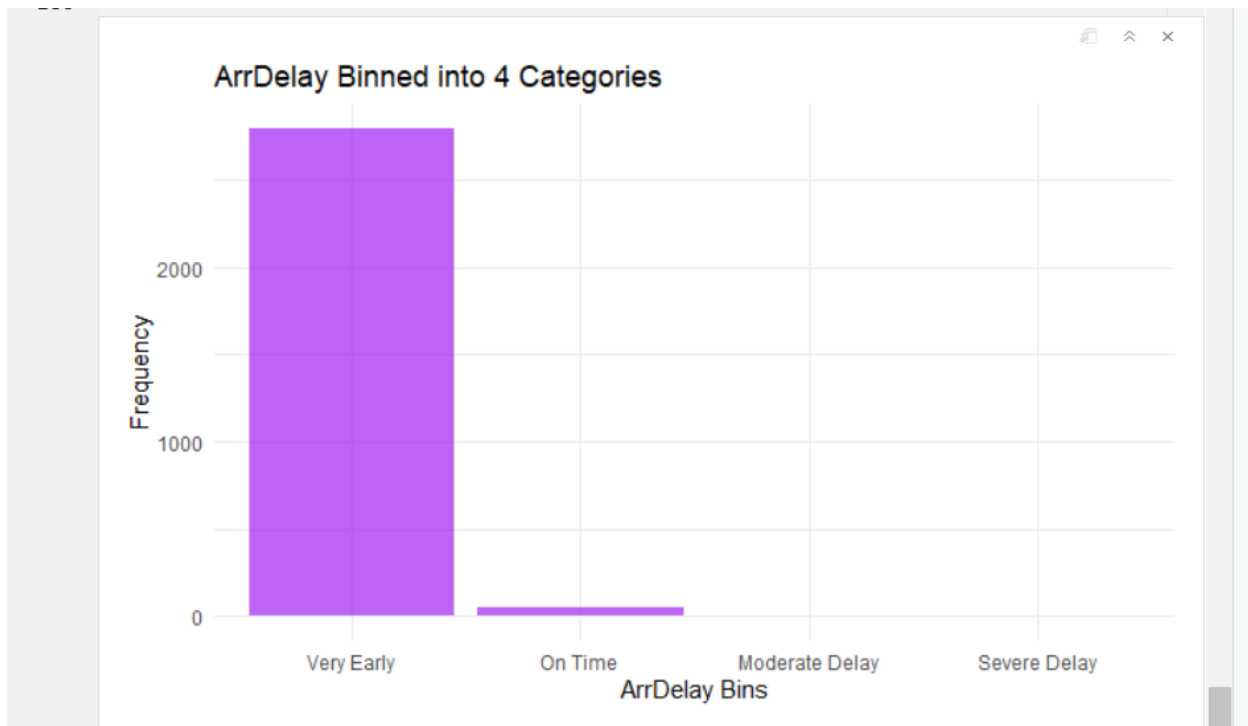
```



```

215 ```{r}
216 # 2. Implement a binning strategy: Divide ArrDelay into 4 bins (quantiles)
217 sub_airline_cleaned <- sub_airline_cleaned %>%
218   mutate(ArrDelay_bins = cut(ArrDelay,
219     breaks = 4, # 4 bins
220     labels = c("Very Early", "On Time", "Moderate Delay",
221       "Severe Delay"),
222     include.lowest = TRUE))
223
224 ```{r}
225 # 3. Visualize the results: Bar plot of the binned data
226 p_bins <- ggplot(sub_airline_cleaned, aes(x = ArrDelay_bins)) +
227   geom_bar(fill = "purple", alpha = 0.7) +
228   ggtitle("ArrDelay Binned into 4 Categories") +
229   xlab("ArrDelay Bins") +
230   ylab("Frequency") +
231   theme_minimal()
232
233 # Print the binned data visualization
234 print(p_bins)
235 ```

```



236 + Insights from the Histogram:

237
238 The histogram shows that most flights have small or moderate delays, with fewer flights experiencing very large delays. There's also a significant portion with early arrivals.

239 + Choice of Bins:

240
241 I chose 4 bins to categorize the flights into meaningful groups: "Very Early",
242 "On Time", "Moderate Delay", and "Severe Delay". This helps simplify the data and makes it easier to analyze delays in broader categories.

243
244 Fewer bins would lose detail, while too many bins might make the data harder to interpret.

245 + Usefulness of Binning:

246
247 Binning is helpful in summarizing continuous data into manageable categories.
248 It can help identify patterns or trends, such as whether flights are mostly on time or delayed, which is useful for reporting and further analysis like building classification models.

249 #TASK 4

250 ```{r}

251 library(dplyr)

252 ```

253

```

254 ~~~{r}
255 # 1. Create dummy variables for the "Reporting_Airline" column
256 # Use model.matrix to create dummy variables for categorical data
257 airline_dummies <- model.matrix(~ Reporting_Airline - 1, data =
  sub_airline_cleaned)
258 airline_dummies <- as.data.frame(airline_dummies) # Convert to data frame
259 ~~~
260
261 ~~~{r}
262 # 2. Create indicator variables for the "Month" column using the "DepDelay" values
263 # Create dummy variables for "Month"
264 month_dummies <- model.matrix(~ Month - 1, data = sub_airline_cleaned)
265 month_dummies <- as.data.frame(month_dummies) # Convert to data frame
266 ~~~
267
268 ~~~{r}
269 # Multiply the dummy variables by "DepDelay" values instead of using 1s and 0s
270 month_dummies_depdelay <- month_dummies * sub_airline_cleaned$DepDelay
271 ~~~
272
273 ~~~{r}
274 # 3. Combine the original data with the new dummy variables
275 sub_airline_transformed <- cbind(sub_airline_cleaned, airline_dummies,
  month_dummies_depdelay)
276 ~~~
277
278 ~~~{r}
279 # Display the first few rows to verify
280 head(sub_airline_transformed)
281 ~~~

```

Description: df [6 × 36]

	Mon...	DayOfWeek	FlightDate	Reporting_Airline	Origin	Dest	CRSDepTime	
	<dbl>	<dbl>	<date>	<chr>	<chr>	<chr>	<chr>	
1	3	5	2003-03-28	UA	LAX	JFK	2210	
2	11	4	2018-11-29	AS	LAX	JFK	1045	
3	8	5	2015-08-28	UA	LAX	JFK	0805	
4	4	7	2003-04-20	DL	LAX	JFK	2205	
5	11	3	2005-11-30	UA	LAX	JFK	0840	
6	4	1	1992-04-06	UA	LAX	JFK	1450	

6 rows | 1-8 of 36 columns

282 + Dataset Structure Change:

283
284 The dataset has more columns now because each category in Reporting_Airline and Month has been split into separate indicator (dummy) columns. For each airline, there's a new column (e.g., Reporting_Airline_AA), and for each month, there's a column multiplied by DepDelay.

285
286 + Benefits:

287
288 Clarity: Indicator variables make it easy to compare the effect of specific airlines or months on delays.

289 Modeling: Many machine learning models (e.g., linear regression, decision trees) work better with numeric data, so converting categorical variables to numeric form (via dummies) makes the data ready for analysis.

290
291 + Drawbacks:

292
293 More Complexity: The dataset grows in size, which might increase computational costs and complexity.

294 Multicollinearity Risk: If too many dummy variables are created without careful consideration, it might introduce multicollinearity issues (redundancy between variables).

295
296 + Usefulness in Analysis:

297
298 These indicator variables will be useful for regression or classification models, allowing the analysis of how specific airlines or months impact delays. They help to capture categorical patterns in a numeric-friendly format for machine learning.

299

300 #TASK 5

301

302 ```{r}

303 # 1. Count the number of data points for each airline

304 airline_counts <- sub_airline_cleaned %>%

305 group_by(Reporting_Airline) %>%

306 summarise(count = n())

307 ```

308

309 ```{r}

310 # 2. Create a bar plot

311 p_airline <- ggplot(airline_counts, aes(x = Reporting_Airline, y = count)) +

312 geom_bar(stat = "identity", fill = "skyblue", color = "black") +

313 ggtitle("Number of Data Points for Each Airline") +

314 xlab("Airline") +

315 ylab("Number of Data Points") +


316 theme_minimal()

317

318 # Print the bar plot

319 print(p_airline)

320 ```



321
322 + Airlines with the Most and Least Data Points:
323
324 The bar plot shows which airlines have the most data points (likely the larger
airlines) and which have the least data points (smaller airlines or those with
fewer flights in the dataset).
325
326 + Impact of Data Point Differences:
327
328 Airlines with more data points will have stronger statistical significance in
analysis, while airlines with fewer data points may provide less reliable results.
This imbalance could skew conclusions if not handled carefully, especially in
modeling or predictions.
329
330 + Additional Visualization Suggestion:
331
332 A box plot showing the distribution of ArrDelay or DepDelay for each airline
could provide insights into how delays vary between airlines, helping to
understand if certain airlines tend to have more delays than others.
333
334